

LECTURE 3:

LOGICAL DATABASE DESIGN AND THE RELATIONAL MODEL

Modern Database Management

OBJECTIVES

- Define terms
- List five properties of relations
- State two properties of candidate keys
- Define first, second, and third normal form
- Describe problems from merging relations
- Transform E-R diagrams to relations
- Create tables with entity and relational integrity constraints
- Use normalization to convert anomalous tables to well-structured relations

COMPONENTS OF RELATIONAL MODEL

- Data structure
 - Tables (relations), rows, columns
- Data manipulation
 - Powerful SQL operations for retrieving and modifying data
- Data integrity
 - Mechanisms for implementing business rules that maintain integrity of manipulated data


RELATION

- A relation is a named, two-dimensional table of data.
- A table consists of rows (records) and columns (attribute or field).
- **Requirements for a table to qualify as a relation:**
 - It must have a unique name.
 - Every attribute value must be atomic (not multivalued, not composite).
 - Every row must be unique (can't have two rows with exactly the same values for all their fields).
 - Attributes (columns) in tables must have unique names.
 - The order of the columns must be irrelevant.
 - The order of the rows must be irrelevant.

CORRESPONDENCE WITH E-R MODEL

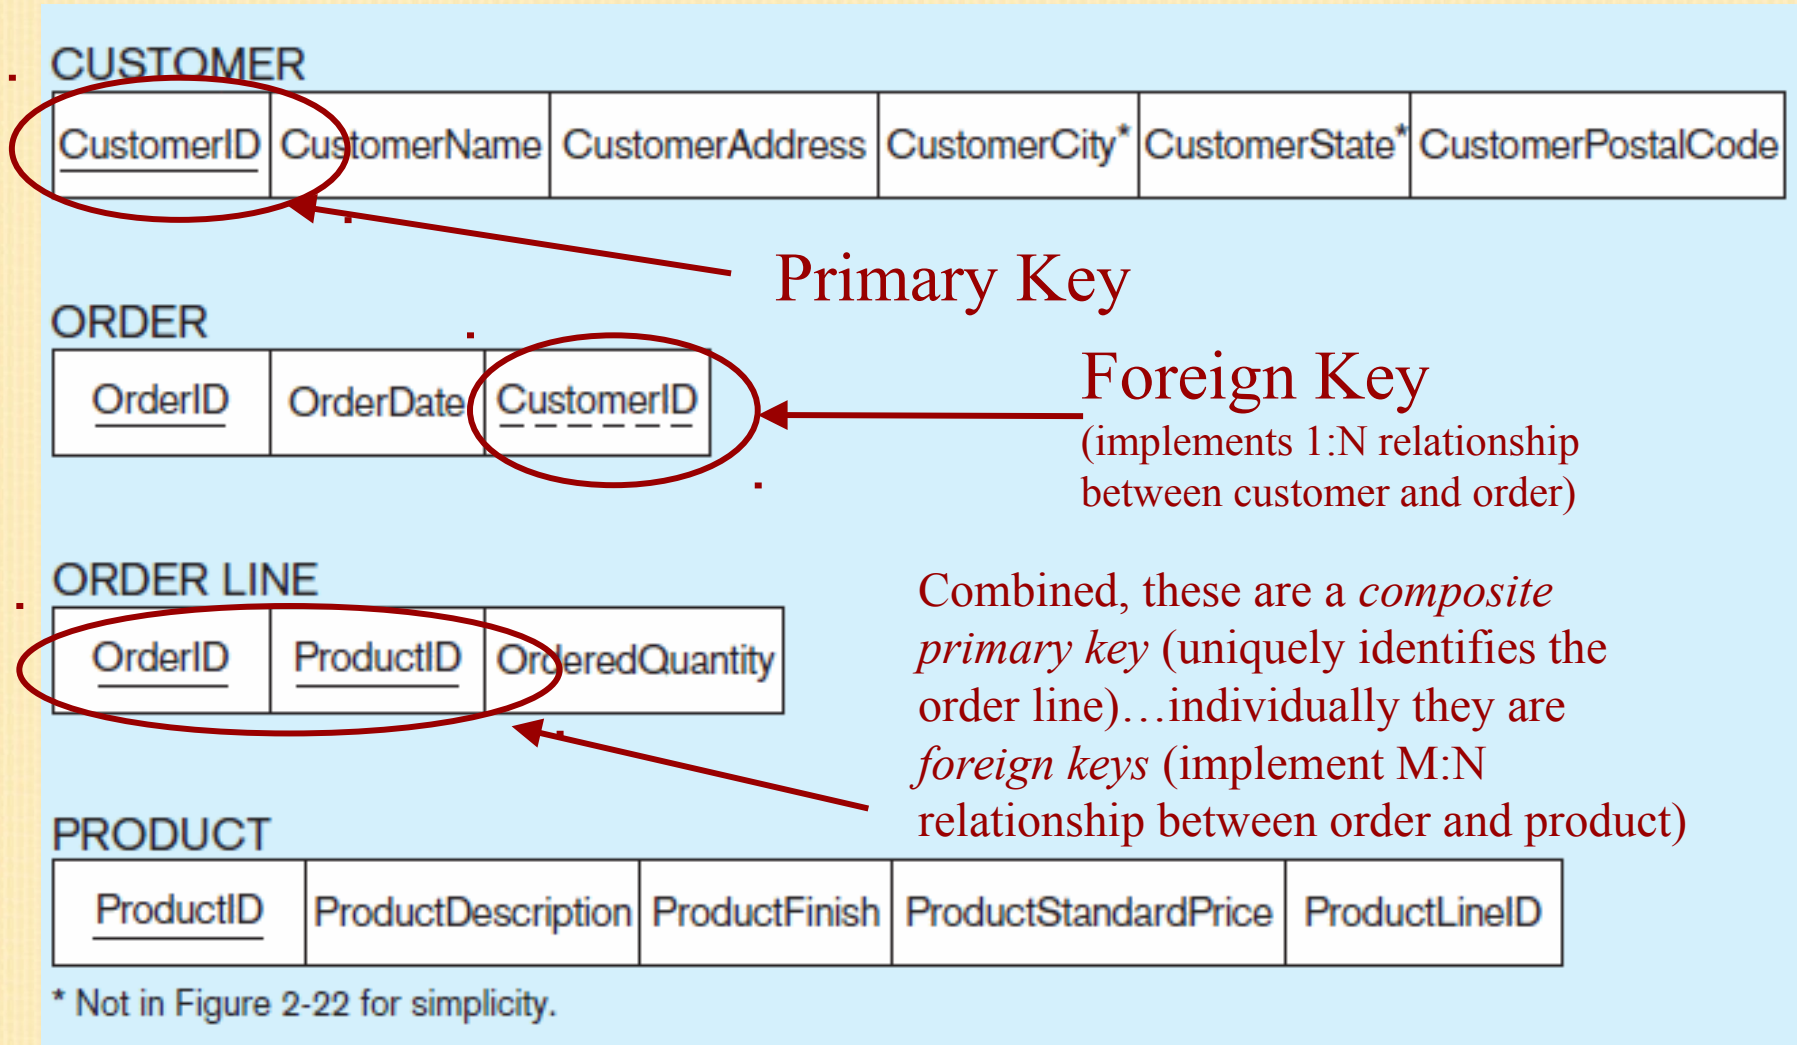
- Relations (tables) correspond with entity types and with many-to-many relationship types.
- Rows correspond with entity instances and with many-to-many relationship instances.
- Columns correspond with attributes.
- NOTE: The word ***relation*** (in a relational database) is NOT the same as the word ***relationship*** (in E-R model).

KEY FIELDS



- Keys are special fields that serve two main purposes:
 - Primary keys** are unique identifiers of the relation. Examples include employee numbers, social security numbers, etc. *This guarantees that all rows are unique.*
 - Foreign keys** are identifiers that enable a dependent relation (on the many side of a relationship) to refer to its parent relation (on the one side of the relationship).
- Keys can be **simple** (a single field) or **composite** (more than one field).
- Keys usually are used as indexes to speed up the response to user queries.

Figure: Example Schema for four relations



INTEGRITY CONSTRAINTS

- Domain Constraints
 - Allowable values for an attribute
- Entity Integrity
 - No primary key attribute may be null.
All primary key fields **MUST** have data.
- Action Assertions
 - Business rules (Recall from Lecture 2)

TABLE 4-1 Domain Definitions for INVOICE Attributes

Attribute	Domain Name	Description	Domain
CustomerID	Customer IDs	Set of all possible customer IDs	character: size 5
CustomerName	Customer Names	Set of all possible customer names	character: size 25
CustomerAddress	Customer Addresses	Set of all possible customer addresses	character: size 30
CustomerCity	Cities	Set of all possible cities	character: size 20
CustomerState	States	Set of all possible states	character: size 2
CustomerPostalCode	Postal Codes	Set of all possible postal zip codes	character: size 10
OrderID	Order IDs	Set of all possible order IDs	character: size 5
OrderDate	Order Dates	Set of all possible order dates	date: format mm/dd/yy
ProductID	Product IDs	Set of all possible product IDs	character: size 5
ProductDescription	Product Descriptions	Set of all possible product descriptions	character: size 25
ProductFinish	Product Finishes	Set of all possible product finishes	character: size 15
ProductStandardPrice	Unit Prices	Set of all possible unit prices	monetary: 6 digits
ProductLineID	Product Line IDs	Set of all possible product line IDs	integer: 3 digits
OrderedQuantity	Quantities	Set of all possible ordered quantities	integer: 3 digits

Domain definitions enforce domain integrity constraints

INTEGRITY CONSTRAINTS

□ **Referential Integrity**—rule states that any foreign key value (on the relation of the many side) **MUST** match a primary key value in the relation of the one side. (Or the foreign key can be null)

□ For example: Delete Rules

□ **Restrict**—don't allow delete of “parent” side if related rows exist in “dependent” side

□ **Cascade**—automatically delete “dependent” side rows that correspond with the “parent” side row to be deleted

□ **Set-to-Null**—set the foreign key in the dependent side to null if deleting from the parent side → not allowed for weak entities

Figure: Referential integrity constraints

CUSTOMER

<u>CustomerID</u>	CustomerName	CustomerAddress	CustomerCity	CustomerState	CustomerPostalCode
-------------------	--------------	-----------------	--------------	---------------	--------------------

ORDER

<u>OrderID</u>	OrderDate	<u>CustomerID</u>
----------------	-----------	-------------------

ORDER LINE

<u>OrderID</u>	<u>ProductID</u>	OrderedQuantity
----------------	------------------	-----------------

PRODUCT

<u>ProductID</u>	ProductDescription	ProductFinish	ProductStandardPrice	ProductLineID
------------------	--------------------	---------------	----------------------	---------------

Referential
integrity
constraints are
drawn via arrows
from dependent to
parent table


TRANSFORMING E-R DIAGRAMS INTO RELATIONS

Mapping Regular Entities to Relations

- **Simple attributes:** E-R attributes map directly onto the relation
- **Composite attributes:** Use only their simple, component attributes
- **Multivalued Attribute:** Becomes a separate relation with a foreign key taken from the superior entity

Figure: Mapping a regular entity

**(a) CUSTOMER
entity type with
simple
attributes**



(b) CUSTOMER relation

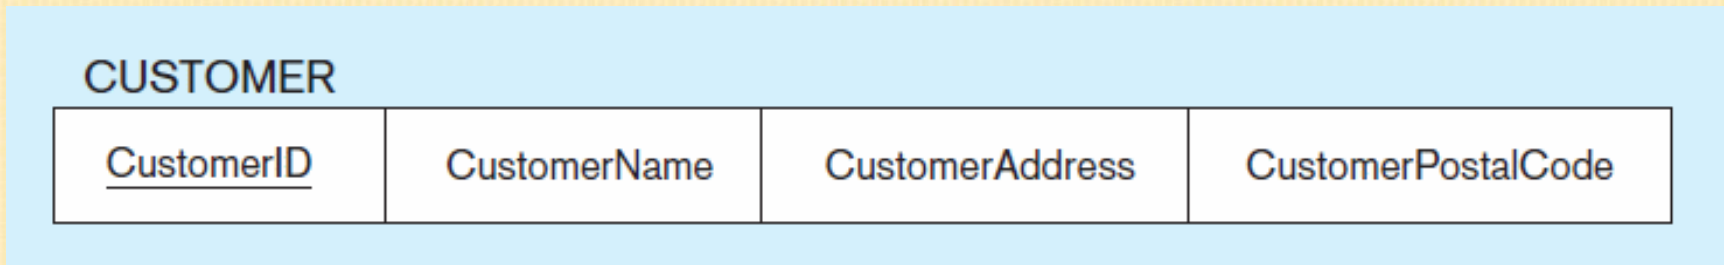
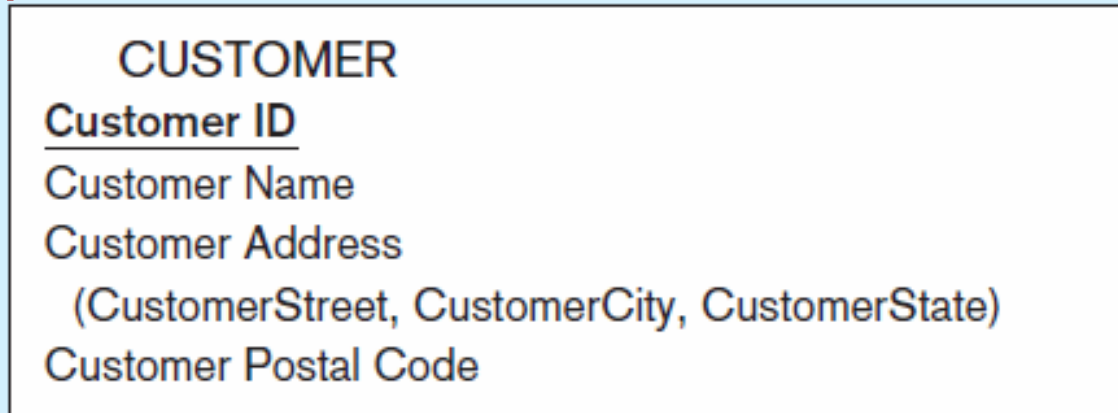


Figure: Mapping a composite attribute

(a) CUSTOMER
entity type with
composite
attribute




(b) CUSTOMER relation with address detail

CUSTOMER

<u>CustomerID</u>	CustomerName	CustomerStreet	CustomerCity	CustomerState	CustomerPostalCode
-------------------	--------------	----------------	--------------	---------------	--------------------

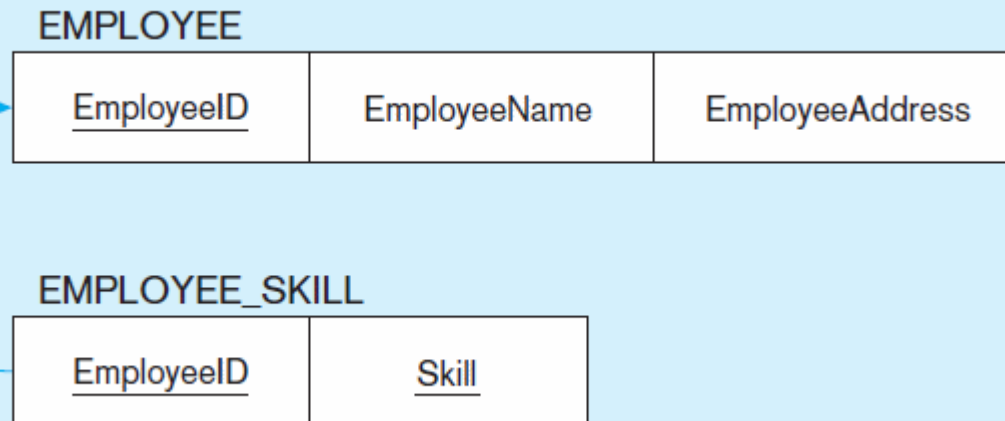
Figure: Mapping an entity with a multivalued attribute

(a)



Multivalued attribute becomes a separate relation with foreign key

(b)



One-to-many relationship between original entity and new relation

TRANSFORMING E-R DIAGRAMS INTO RELATIONS (CONT.)

Mapping Weak Entities

- ▮ Becomes a separate relation with a foreign key taken from the superior entity
- ▮ Primary key composed of:
 - ▮ Partial identifier of weak entity
 - ▮ Primary key of identifying relation (strong entity)

Figure: Example of mapping a weak entity

a) Weak entity DEPENDENT

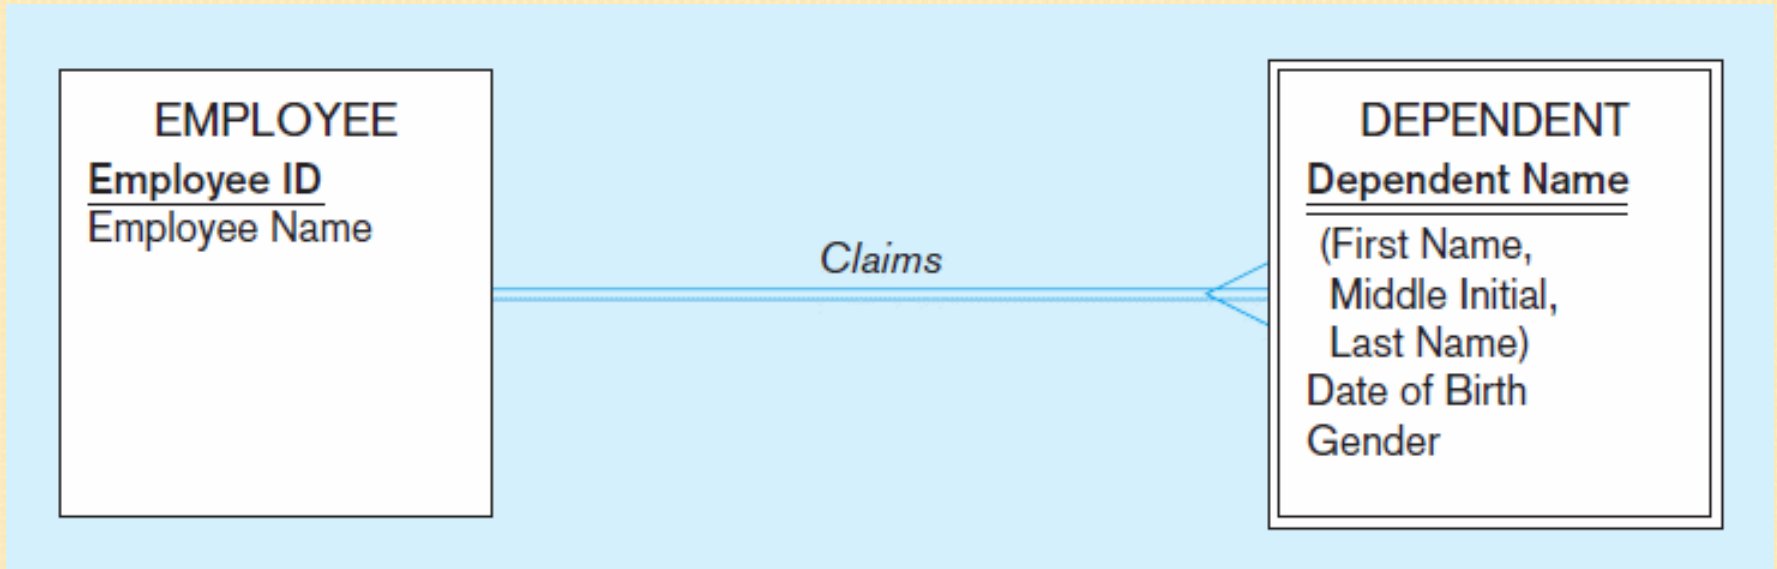
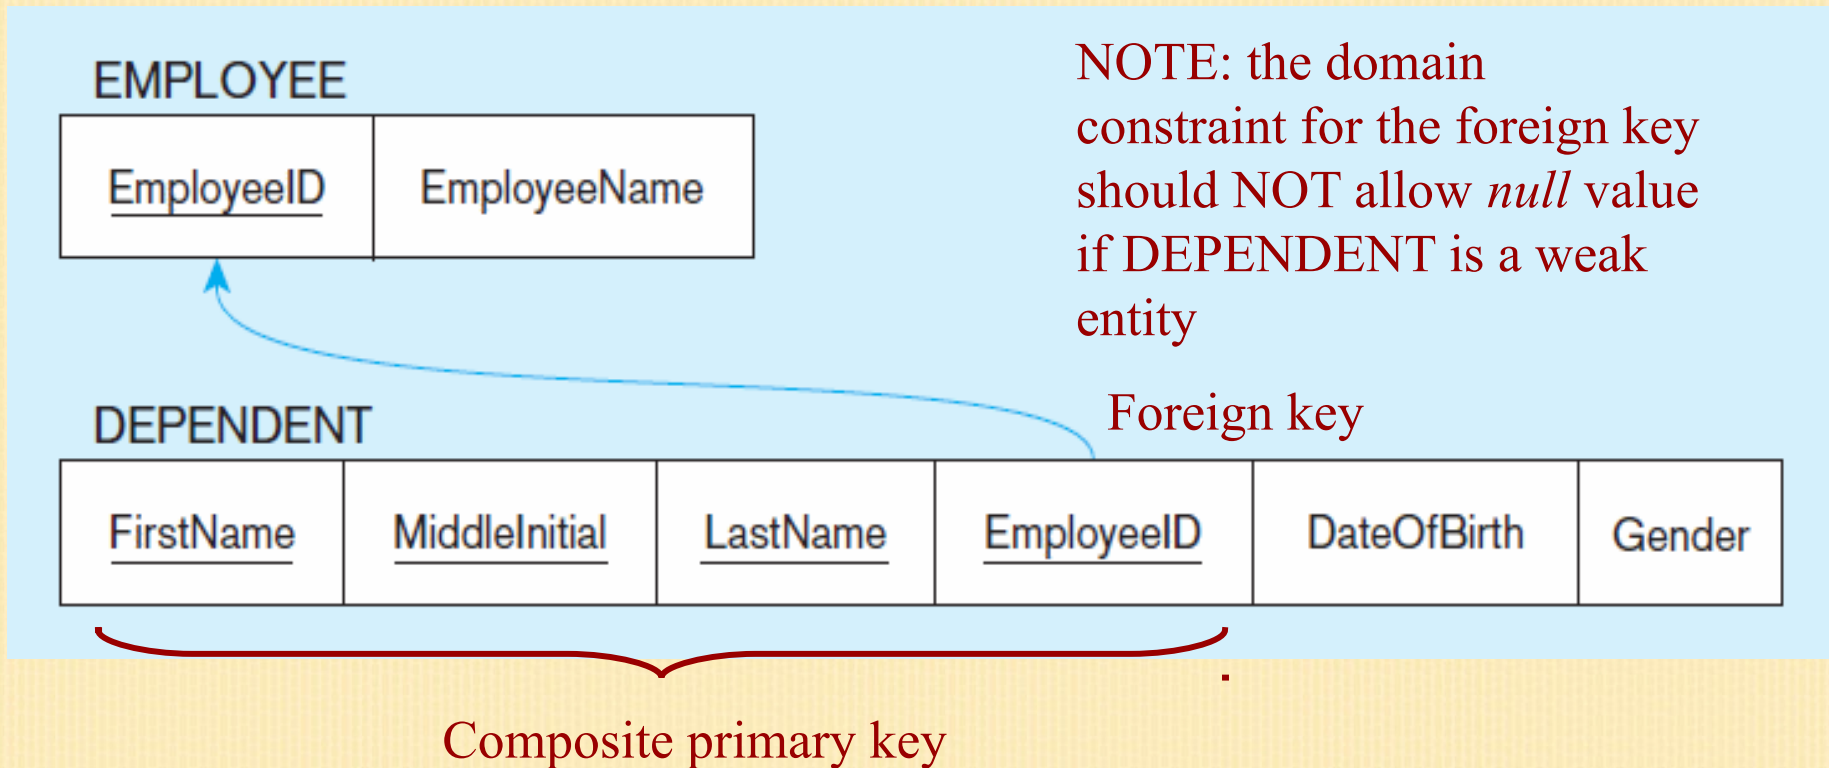


Figure: Example of mapping a weak entity (cont.)

b) Relations resulting from weak entity



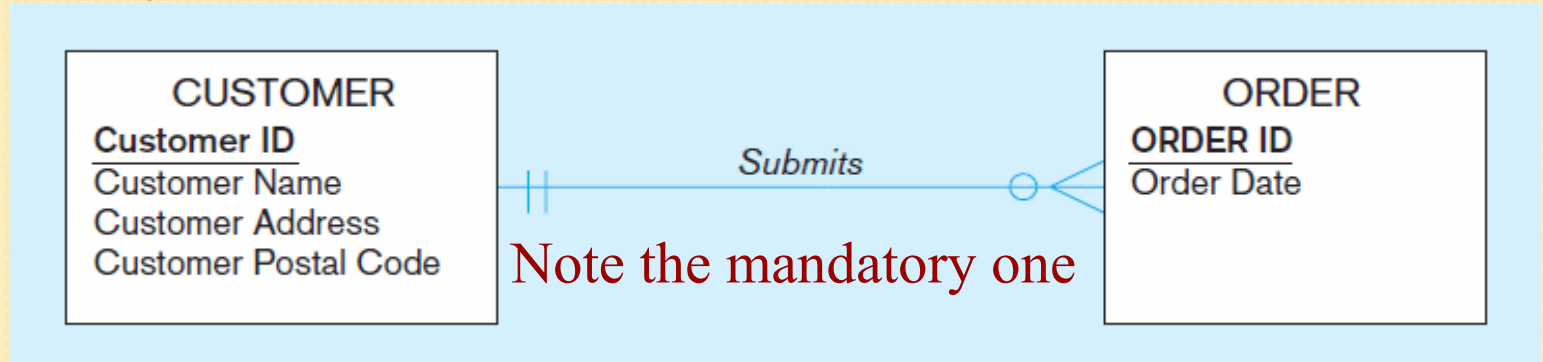
TRANSFORMING E-R DIAGRAMS INTO RELATIONS (CONT.)

Mapping Binary Relationships

- **One-to-Many**—Primary key on the one side becomes a foreign key on the many side
- **Many-to-Many**—Create a ***new relation*** with the primary keys of the two entities as its primary key
- **One-to-One**—Primary key on mandatory side becomes a foreign key on optional side

Figure: Example of mapping a 1:M relationship

a) Relationship between customers and orders



b) Mapping the relationship

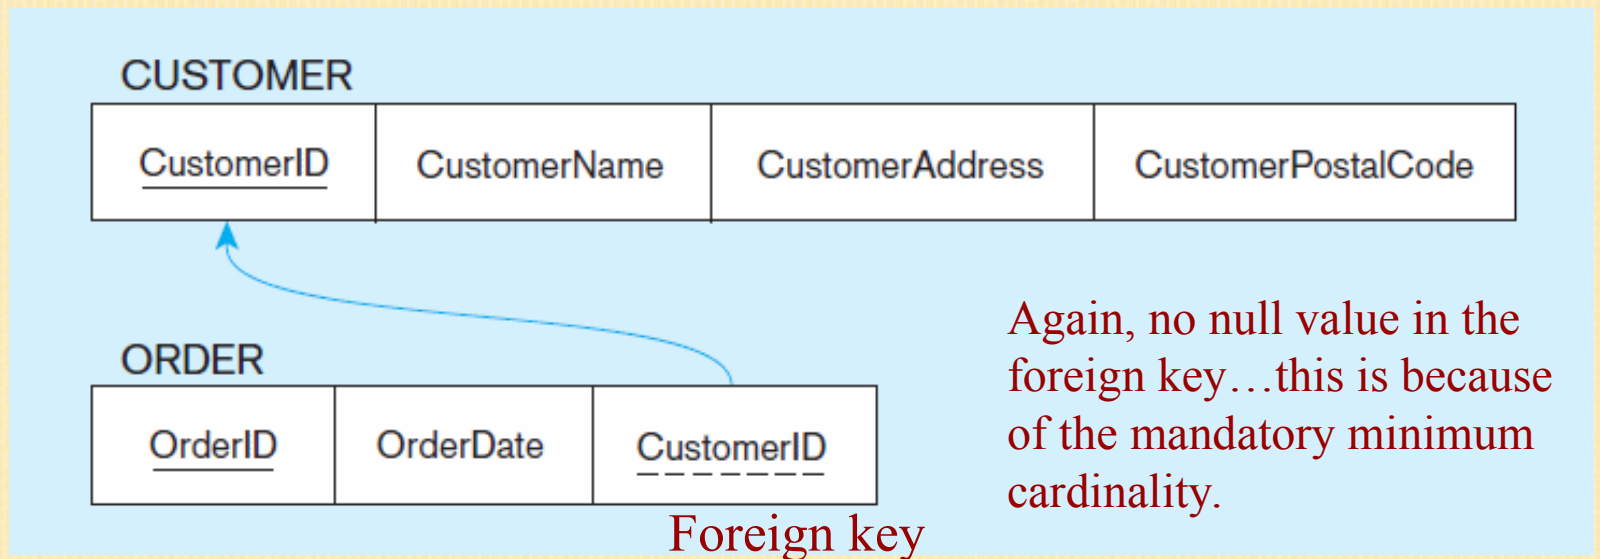
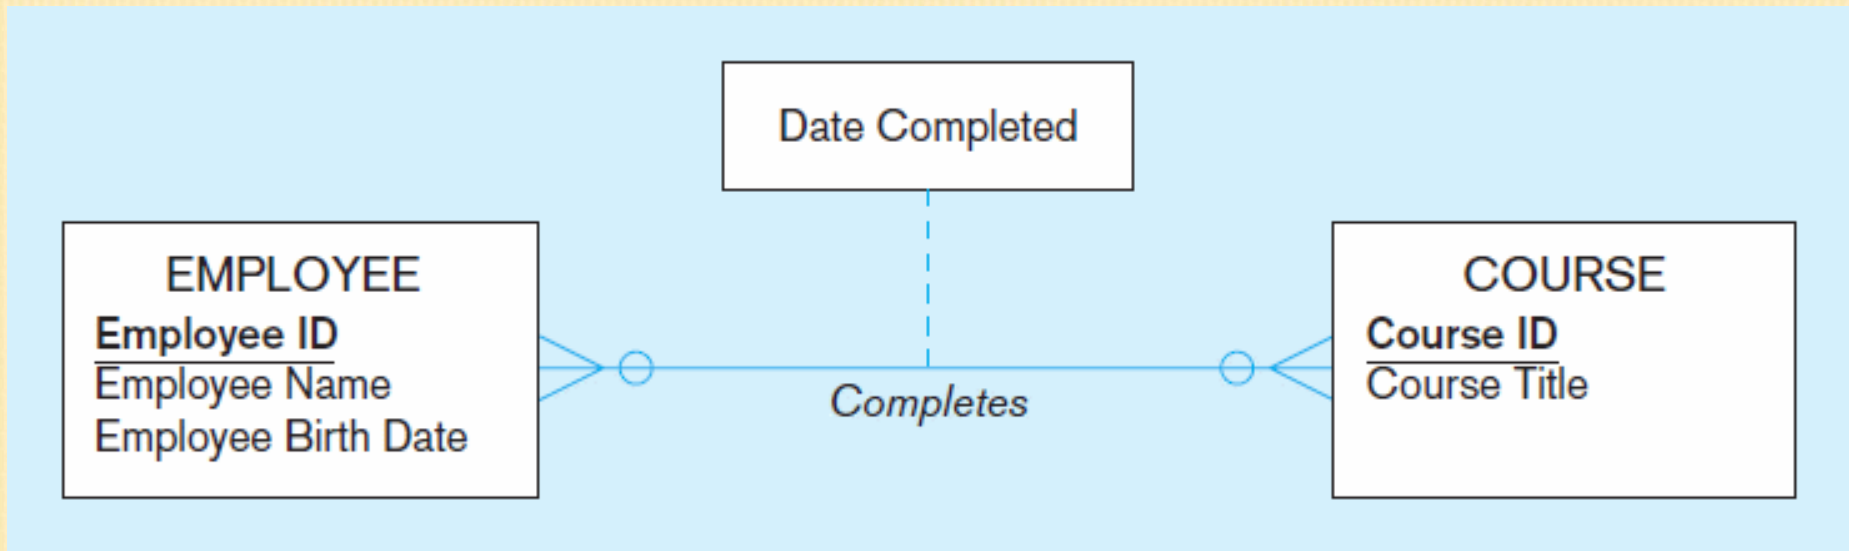


Figure: Example of mapping an M:N relationship

a) Completes relationship (M:N)



The *Completes* relationship will need to become a separate relation.

Figure: Example of mapping an M:N relationship (cont.)

b) Three resulting relations

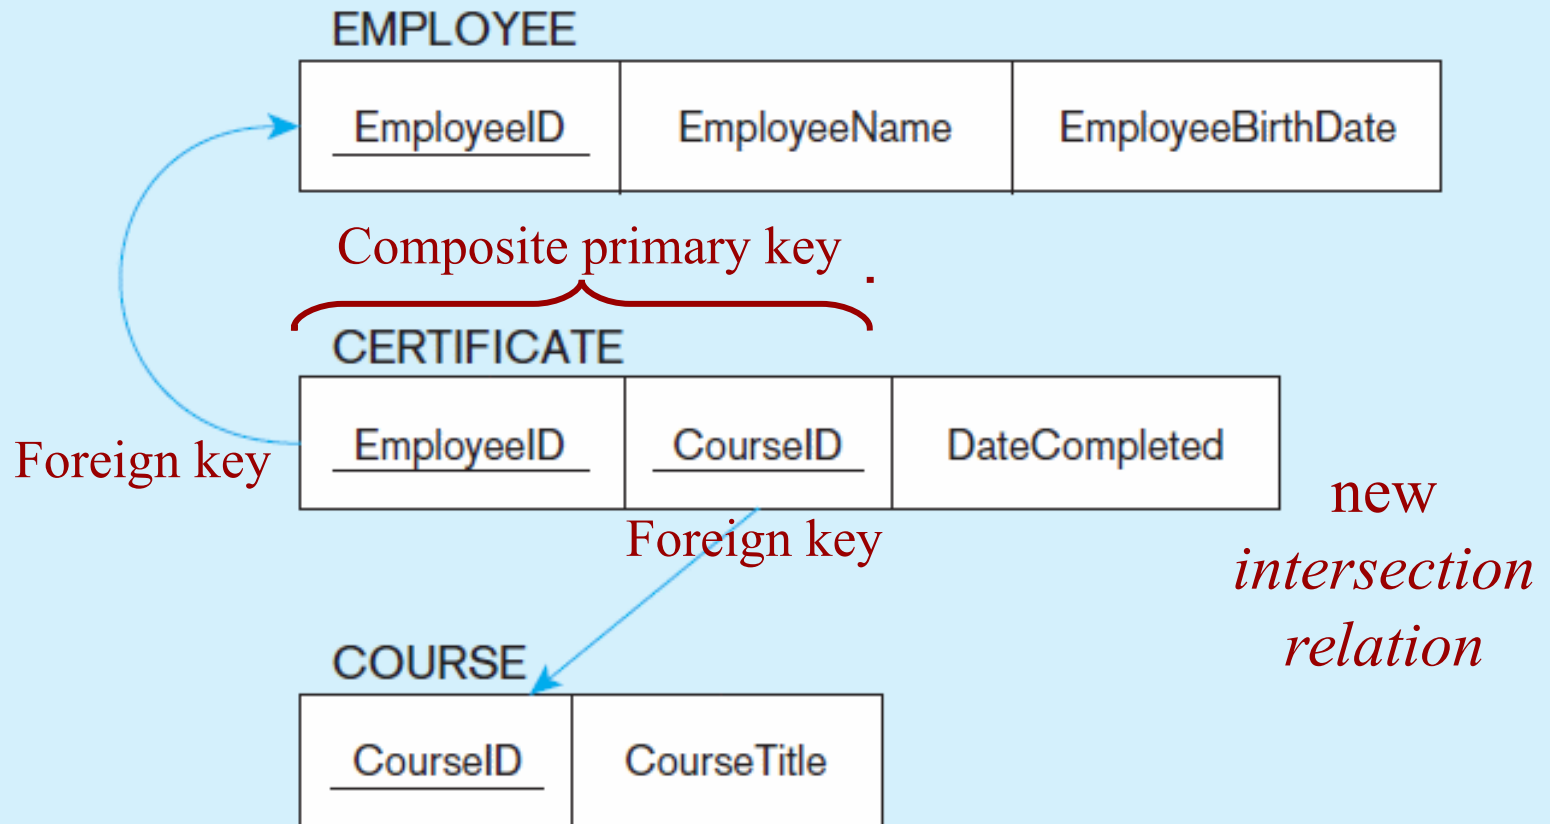
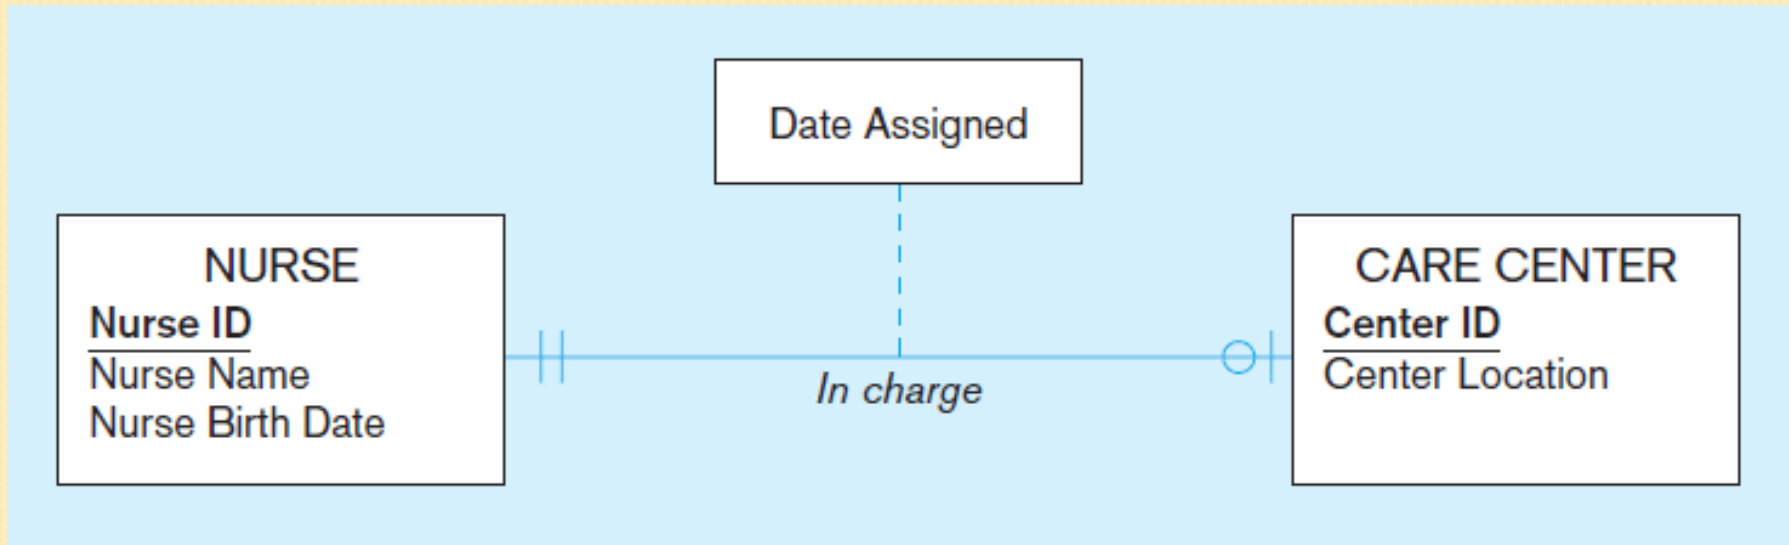


Figure: Example of mapping a binary 1:1 relationship

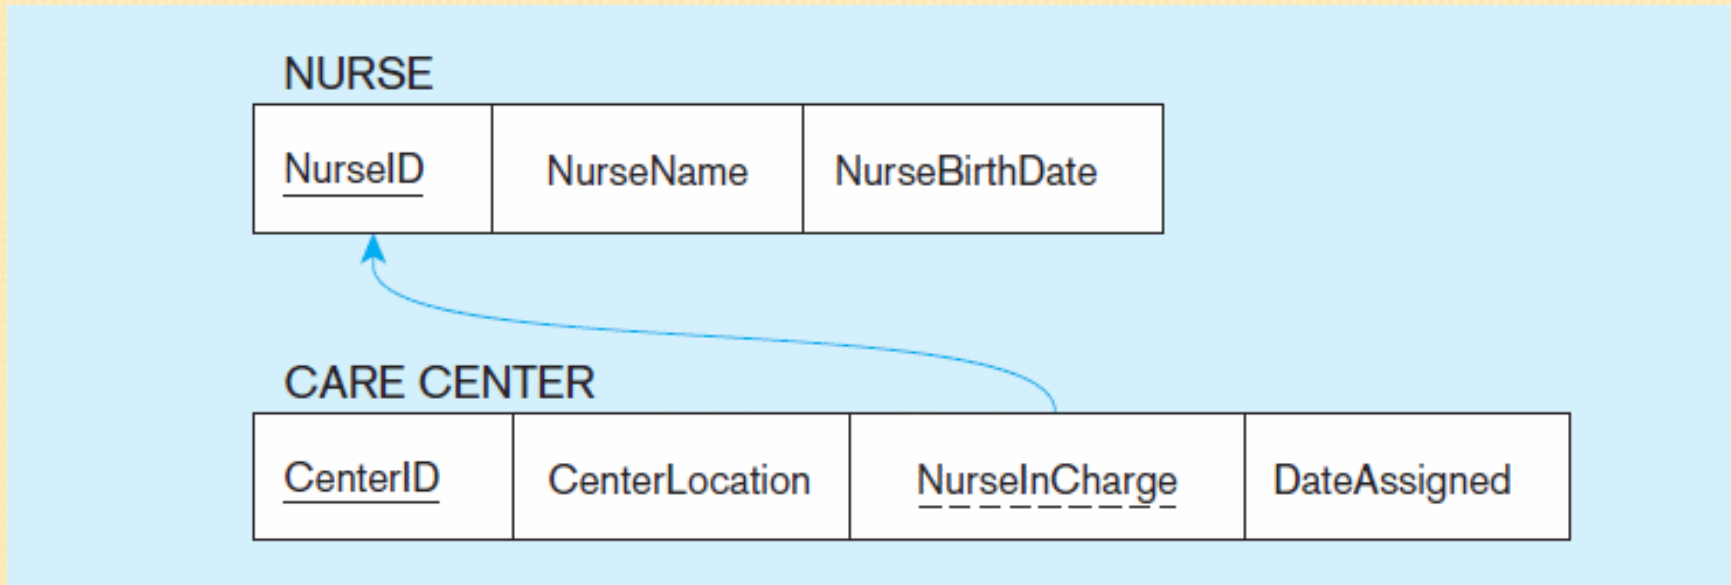
a) In charge relationship (1:1)



Often in 1:1 relationships, one direction is optional

Figure: Example of mapping a binary 1:1 relationship (cont.)

b) Resulting relations



Foreign key goes in the relation on the optional side, matching the primary key on the mandatory side

TRANSFORMING E-R DIAGRAMS INTO RELATIONS (CONT.)

Mapping Associative Entities

- Identifier Not Assigned
 - Default primary key for the association relation is composed of the primary keys of the two entities (as in M:N relationship)
- Identifier Assigned
 - It is natural and familiar to end-users
 - Default identifier may not be unique

Figure: Example of mapping an associative entity

a) An associative entity

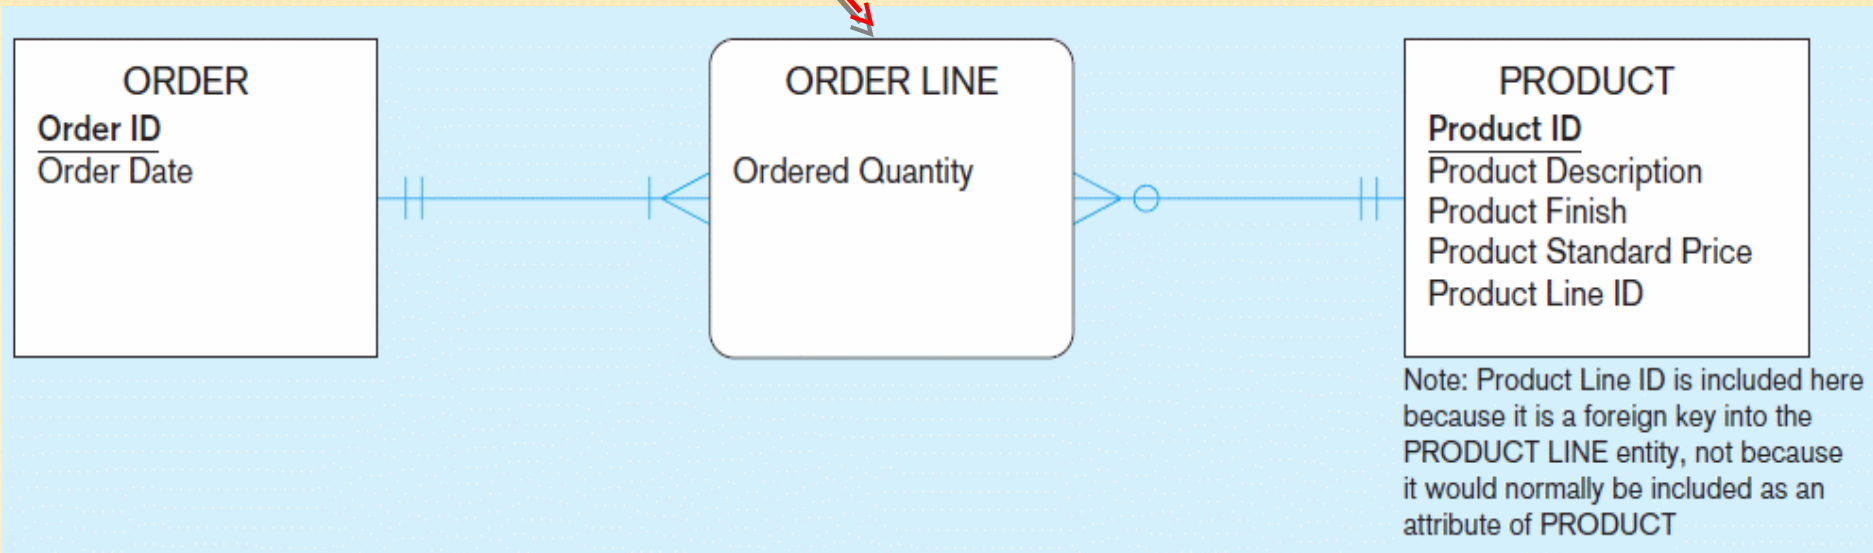
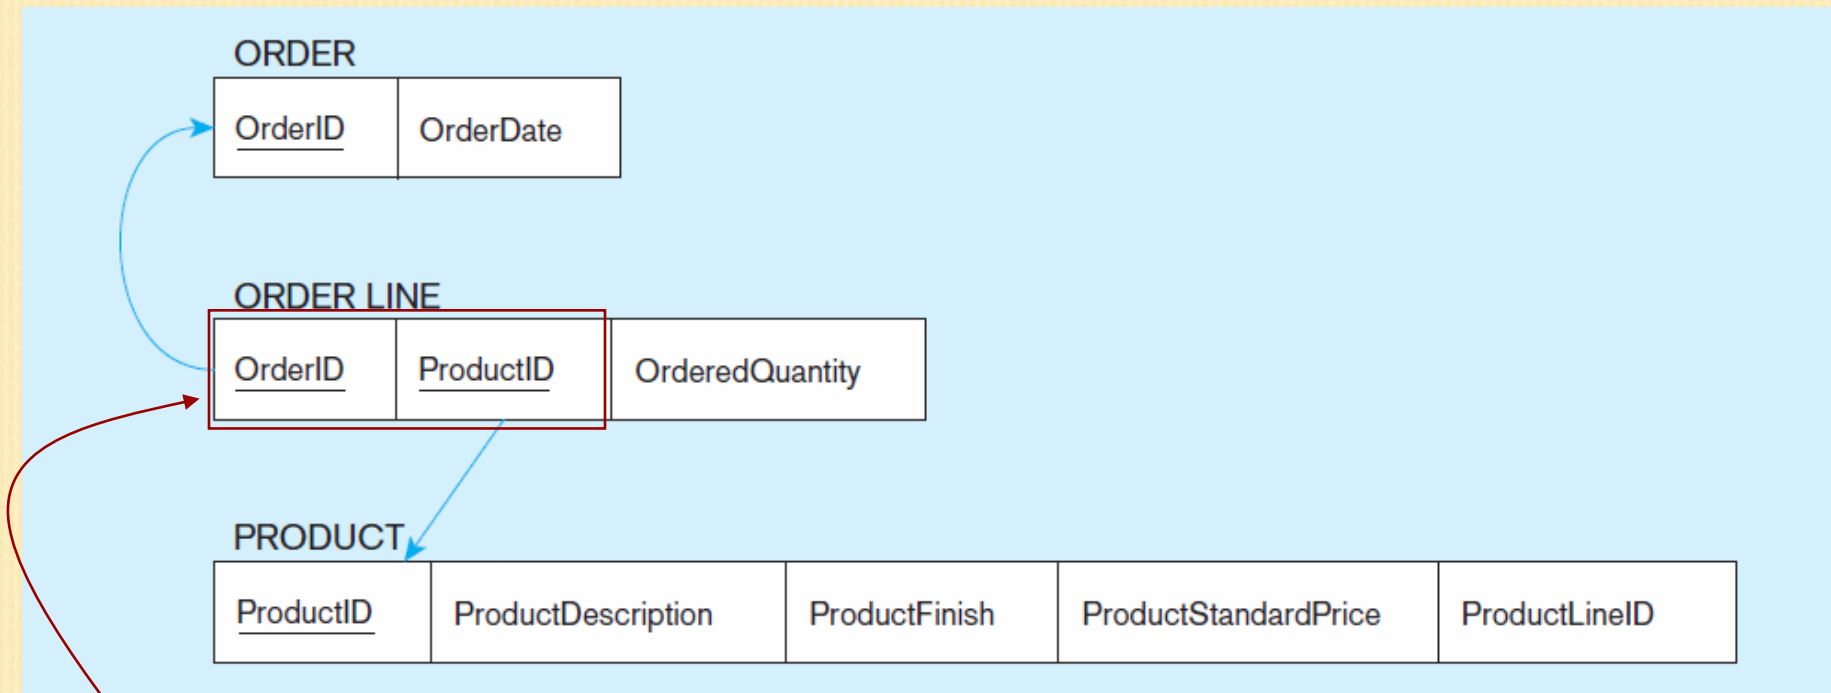


Figure: Example of mapping an associative entity (cont.)

b) Three resulting relations



Composite primary key formed from the two foreign keys

Figure: Example of mapping an associative entity with an identifier

a) SHIPMENT associative entity

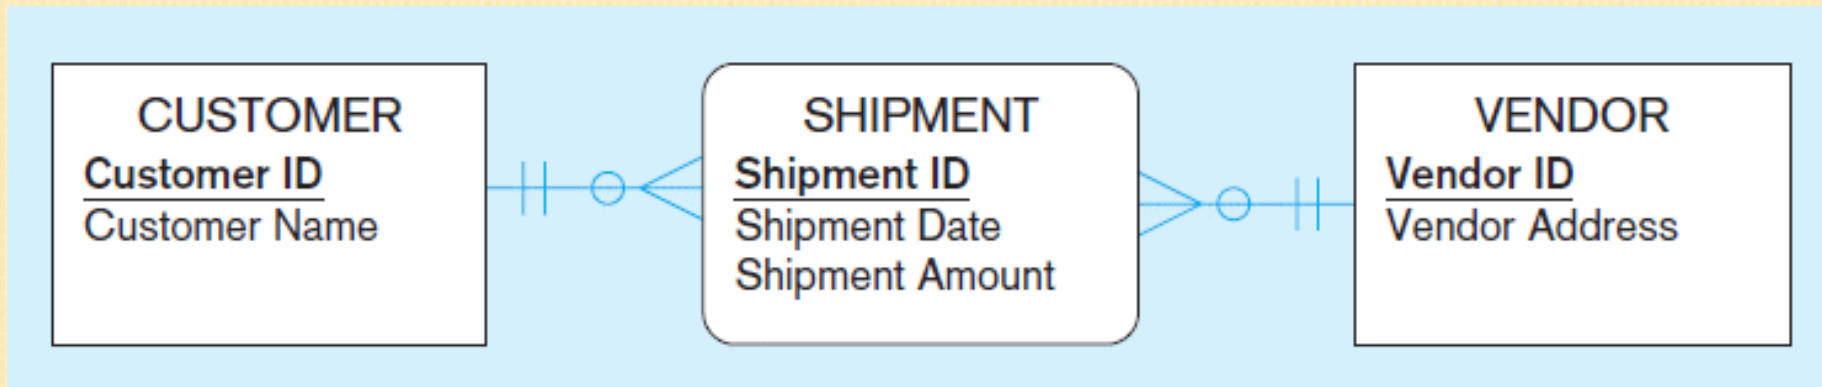
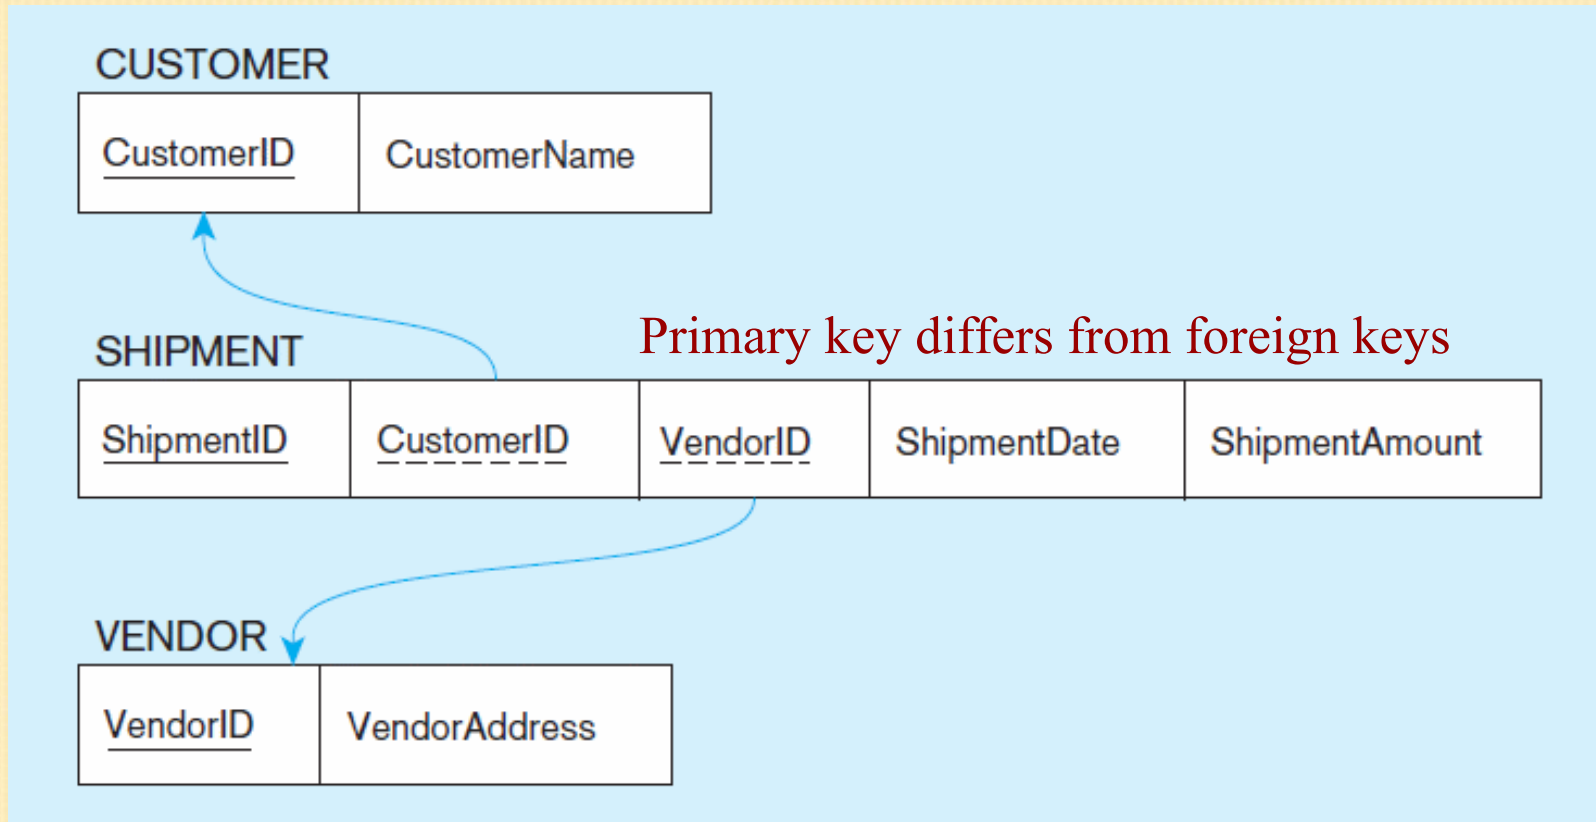


Figure: Example of mapping an associative entity with an identifier (cont.)

b) Three resulting relations



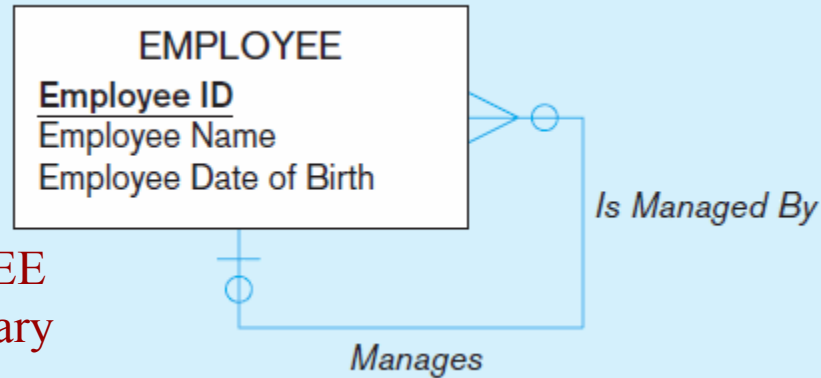
TRANSFORMING E-R DIAGRAMS INTO RELATIONS (CONT.)

Mapping Unary Relationships

- One-to-Many-Recursive foreign key in the same relation
- Many-to-Many-Two relations:
 - One for the entity type
 - One for an associative relation in which the primary key has two attributes, both taken from the primary key of the entity

Figure: Mapping a unary 1:N relationship

(a) EMPLOYEE
entity with unary
relationship



(b)
EMPLOYEE
relation with
recursive
foreign key


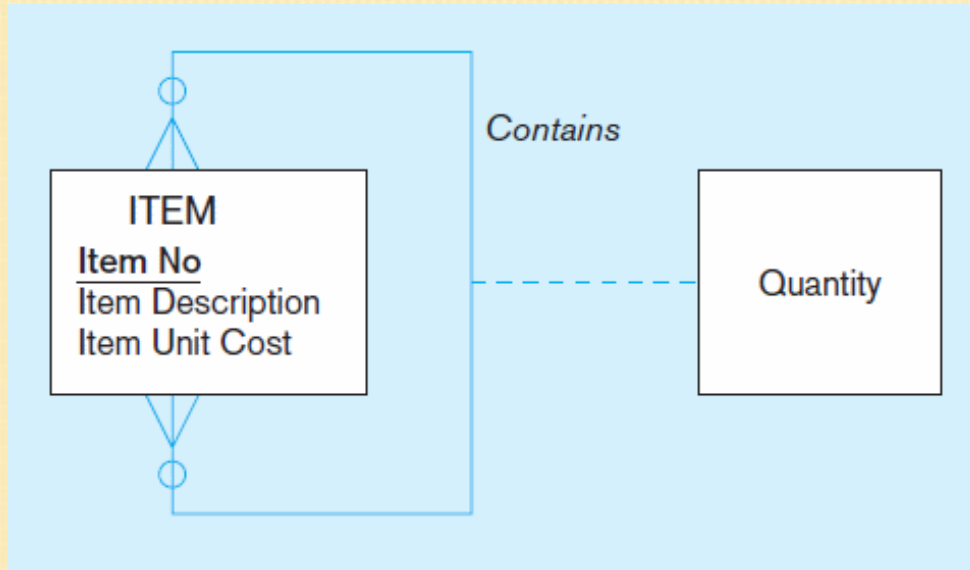
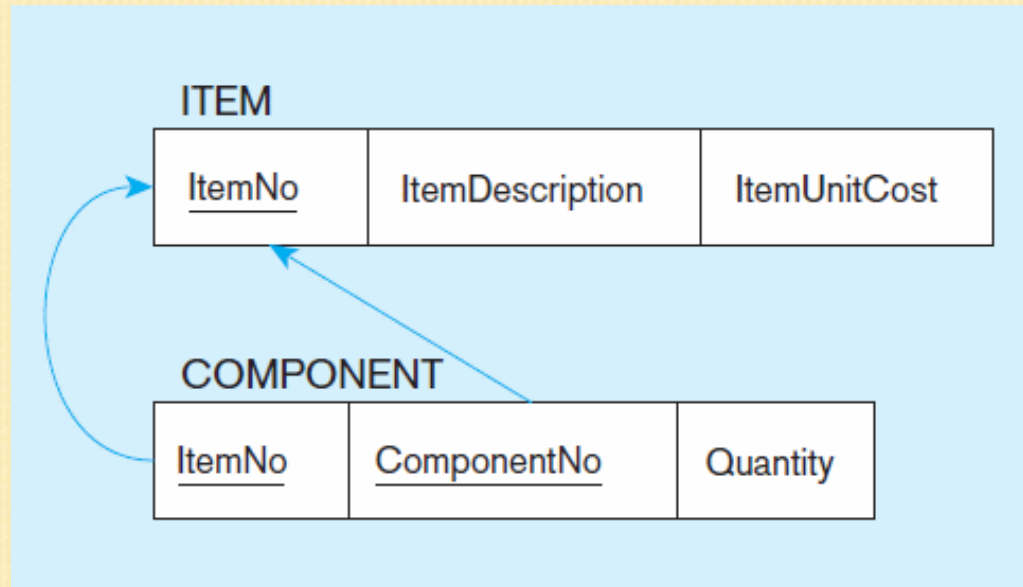


Figure: Mapping a unary M:N relationship



(a) Bill-of-materials relationships (M:N)

(b) ITEM and COMPONENT relations



TRANSFORMING E-R DIAGRAMS INTO RELATIONS (CONT.)

Mapping Ternary (and n-ary) Relationships

- One relation for each entity and one for the associative entity
- Associative entity has foreign keys to each entity in the relationship

Figure: Mapping a ternary relationship

a) PATIENT TREATMENT Ternary relationship with associative entity

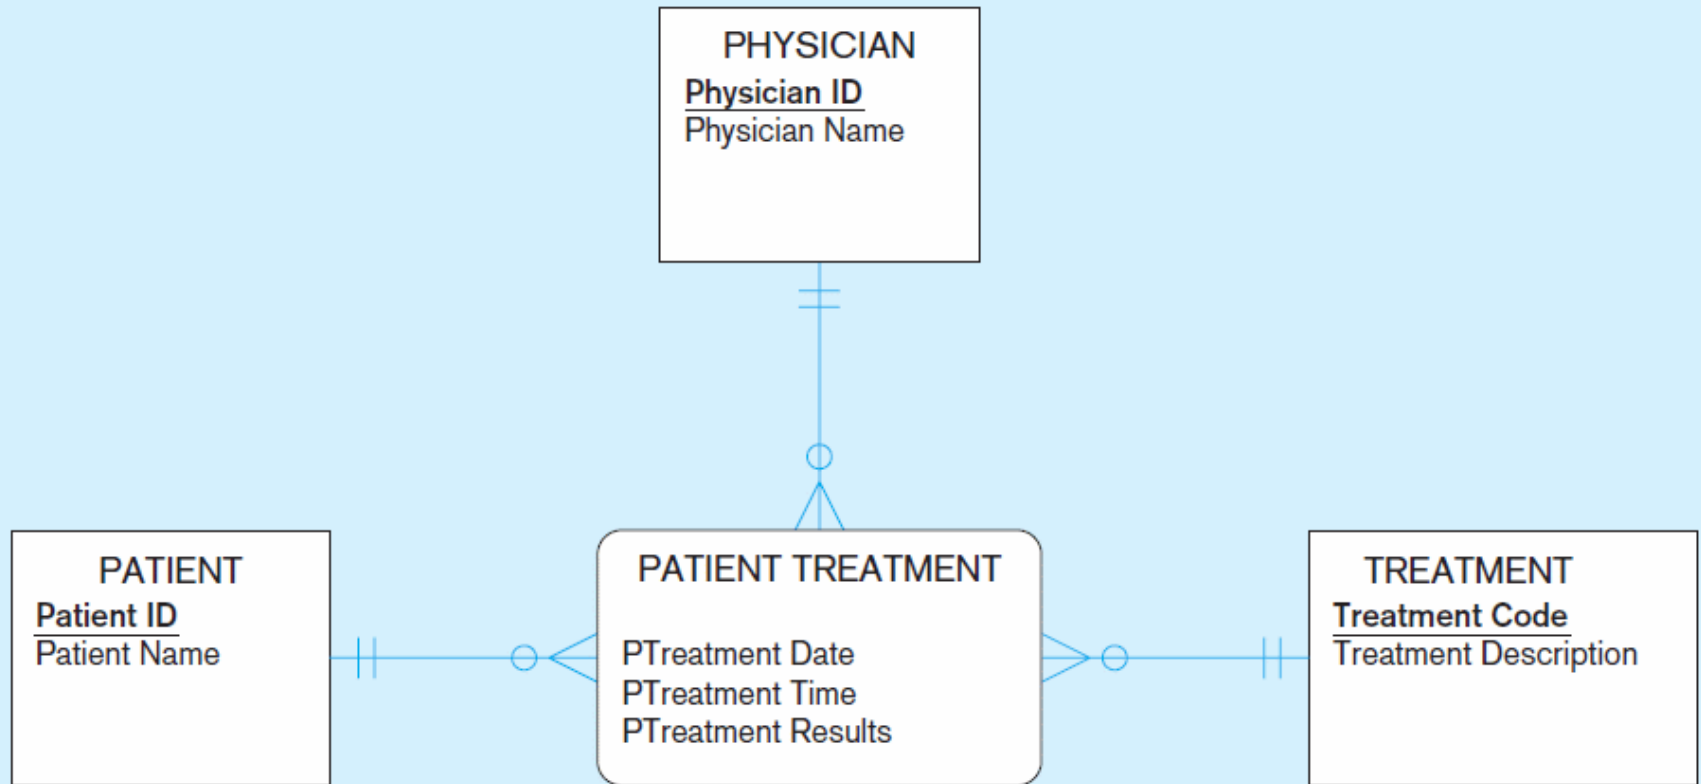
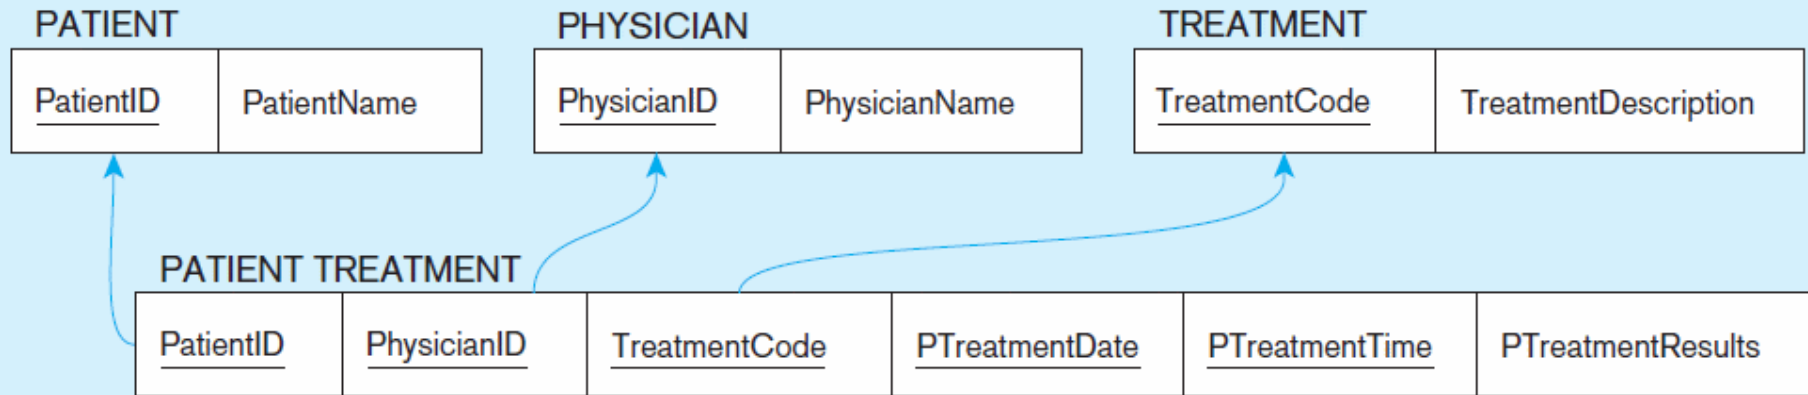


Figure: Mapping a ternary relationship (cont.)

b) Mapping the ternary relationship PATIENT TREATMENT



Remember that the primary key **MUST** be unique.

This is why treatment date and time are included in the composite primary key.

But this makes a very cumbersome key...

It would be better to create a surrogate key like Treatment#.

DATA NORMALIZATION

- Primarily a tool to validate and improve a logical design so that it satisfies certain constraints that ***avoid unnecessary duplication of data***
- The process of decomposing relations with anomalies to produce smaller, ***well-structured*** relations

WELL-STRUCTURED RELATIONS

- A relation that contains minimal data redundancy and allows users to insert, delete, and update rows without causing data inconsistencies
- Goal is to avoid anomalies
 - **Insertion Anomaly**—adding new rows forces user to create duplicate data
 - **Deletion Anomaly**—deleting rows may cause a loss of data that would be needed for other future rows
 - **Modification Anomaly**—changing data in a row forces changes to other rows because of duplication

General rule of thumb: A table should not pertain to more than one entity type.

EXAMPLE

EMPLOYEE2

<u>EmplID</u>	Name	DeptName	Salary	CourseTitle	DateCompleted
100	Margaret Simpson	Marketing	48,000	SPSS	6/19/201X
100	Margaret Simpson	Marketing	48,000	Surveys	10/7/201X
140	Alan Beeton	Accounting	52,000	Tax Acc	12/8/201X
110	Chris Lucero	Info Systems	43,000	Visual Basic	1/12/201X
110	Chris Lucero	Info Systems	43,000	C++	4/22/201X
190	Lorenzo Davis	Finance	55,000		
150	Susan Martin	Marketing	42,000	SPSS	6/19/201X
150	Susan Martin	Marketing	42,000	Java	8/12/201X

Question—Is this a relation?

Answer—Yes: Unique rows and no multivalued attributes

Question—What's the primary key?

Answer—Composite: EmpID, CourseTitle

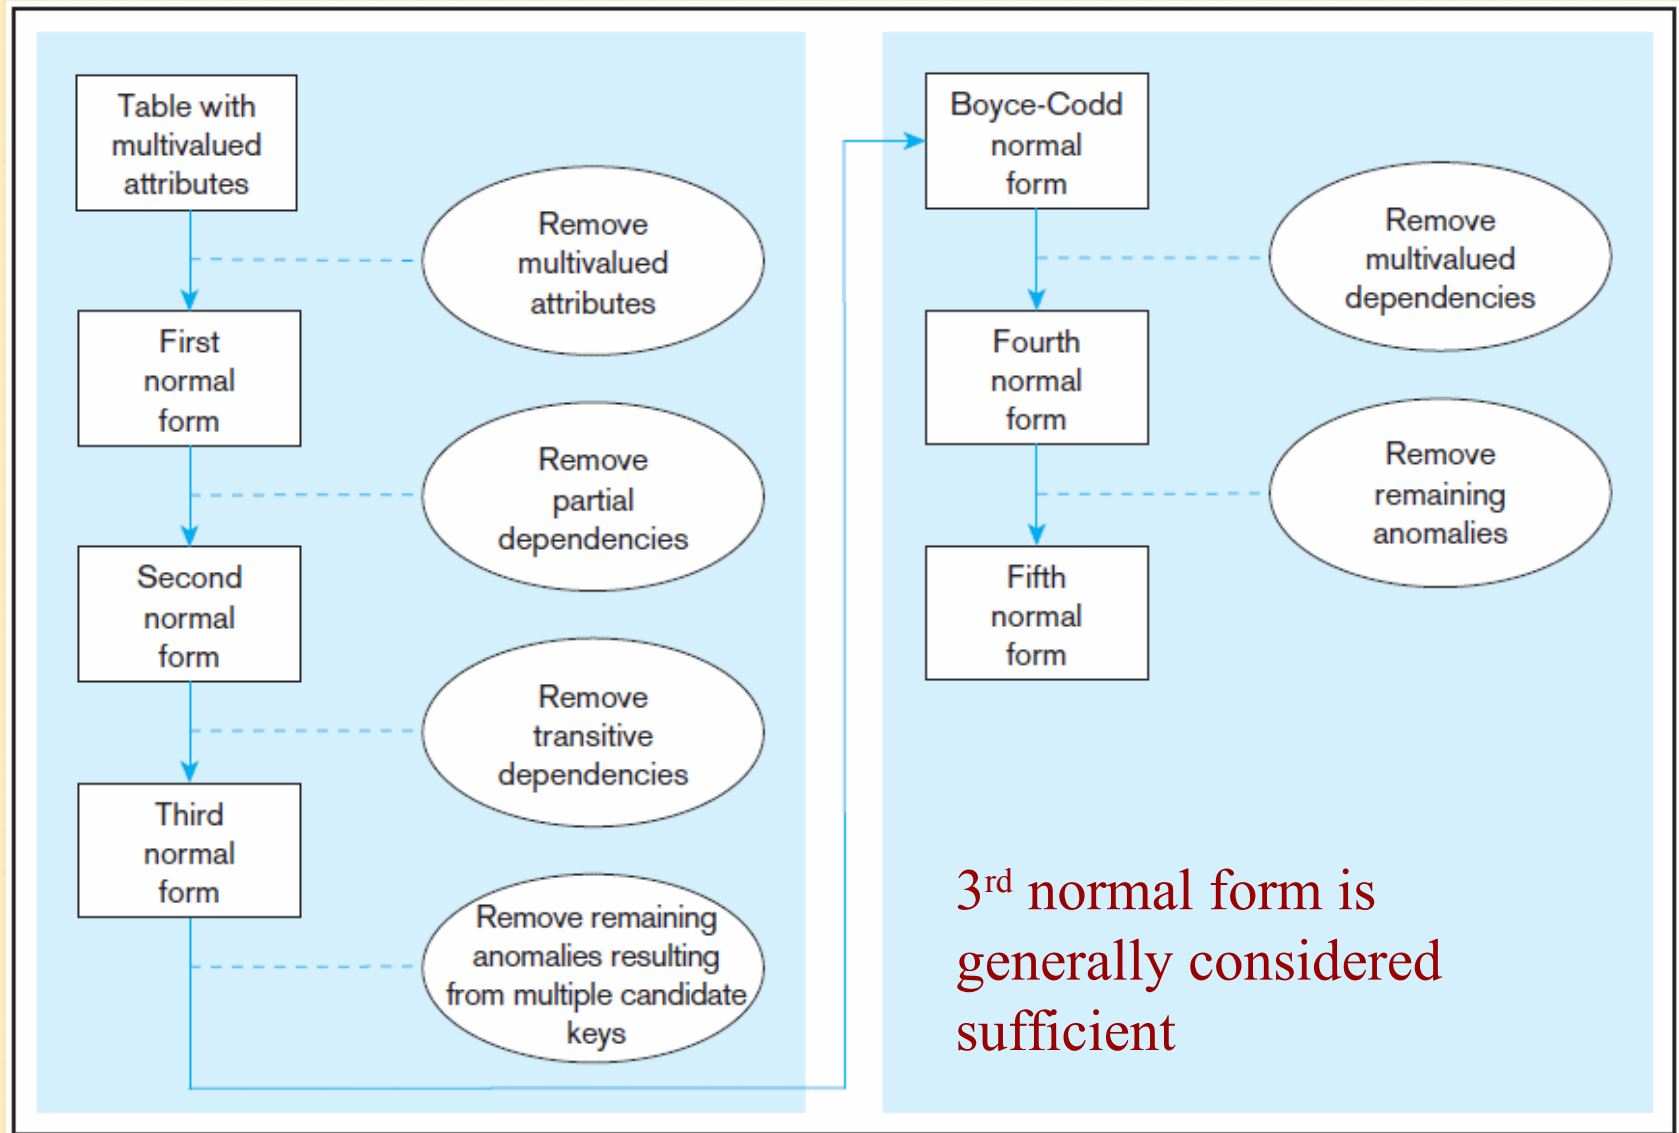
ANOMALIES IN THIS TABLE

- **Insertion**—can't enter a new employee without having the employee take a class (or at least empty fields of class information)
- **Deletion**—if we remove employee 140, we lose information about the existence of a Tax Acc class
- **Modification**—giving a salary increase to employee 100 forces us to update multiple records

Why do these anomalies exist?

Because there are two themes (entity types) in this one relation. This results in data duplication and an unnecessary dependency between the entities.

Figure: Steps in normalization



FUNCTIONAL DEPENDENCIES AND KEYS

- Functional Dependency: The value of one attribute (the **determinant**) determines the value of another attribute
- Candidate Key:
 - A unique identifier. One of the candidate keys will become the primary key
 - E.g., perhaps there is both credit card number and SS# in a table...in this case both are candidate keys.
 - Each non-key field is functionally dependent on every candidate key.

FIRST NORMAL FORM

- No multivalued attributes
- Every attribute value is atomic
- ***All relations are in 1st Normal Form.***

Table with multivalued attributes, not in 1st normal form

<u>OrderID</u>	Order Date	Customer ID	Customer Name	Customer Address	<u>ProductID</u>	Product Description	Product Finish	Product StandardPrice	Ordered Quantity
1006	10/24/2010	2	Value Furniture	Plano, TX	7	Dining Table	Natural Ash	800.00	2
					5	Writer's Desk	Cherry	325.00	2
					4	Entertainment Center	Natural Maple	650.00	1
1007	10/25/2010	6	Furniture Gallery	Boulder, CO	11	4-Dr Dresser	Oak	500.00	4
					4	Entertainment Center	Natural Maple	650.00	3

FIGURE 4-25 INVOICE data (Pine Valley Furniture Company)

Note: This is NOT a relation.

Table with no multivalued attributes and unique rows, in 1st normal form

<u>OrderID</u>	Order Date	Customer ID	Customer Name	Customer Address	<u>ProductID</u>	Product Description	Product Finish	Product StandardPrice	Ordered Quantity
1006	10/24/2010	2	Value Furniture	Plano, TX	7	Dining Table	Natural Ash	800.00	2
1006	10/24/2010	2	Value Furniture	Plano, TX	5	Writer's Desk	Cherry	325.00	2
1006	10/24/2010	2	Value Furniture	Plano, TX	4	Entertainment Center	Natural Maple	650.00	1
1007	10/25/2010	6	Furniture Gallery	Boulder, CO	11	4-Dr Dresser	Oak	500.00	4
1007	10/25/2010	6	Furniture Gallery	Boulder, CO	4	Entertainment Center	Natural Maple	650.00	3

FIGURE 4-26 INVOICE relation (1NF) (Pine Valley Furniture Company)

Note: This is a relation, but not a well-structured one.

ANOMALIES IN THIS TABLE

- **Insertion**—if new product is ordered for order 1007 of existing customer, customer data must be re-entered, causing duplication
- **Deletion**—if we delete the Dining Table from Order 1006, we lose information concerning this item's finish and price
- **Update**—changing the price of product ID 4 requires update in multiple records

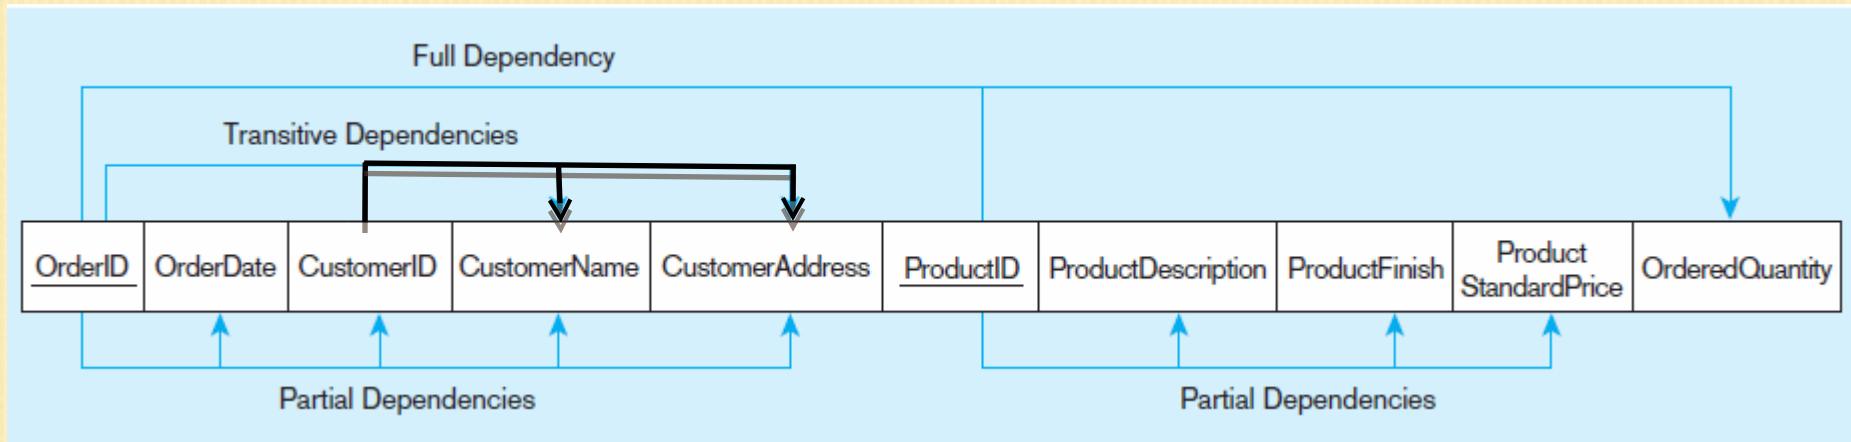
Why do these anomalies exist?

Because there are multiple themes (entity types) in one relation. This results in duplication and an unnecessary dependency between the entities.

SECOND NORMAL FORM

- 1NF PLUS ***every non-key attribute is fully functionally dependent on the ENTIRE primary key***
 - Every non-key attribute must be defined by the entire key, not by only part of the key
 - No partial functional dependencies

Functional dependency diagram for INVOICE



OrderID → OrderDate, CustomerID, CustomerName, CustomerAddress

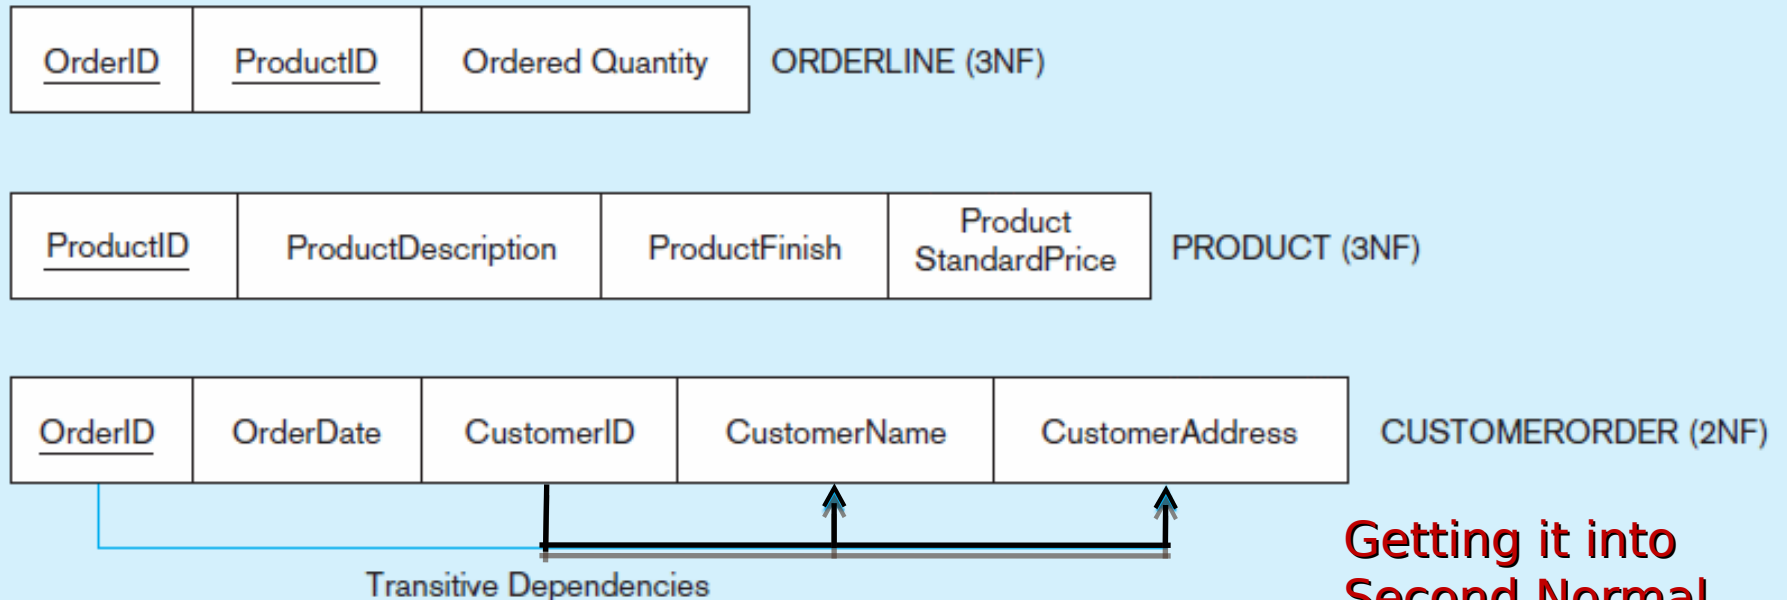
CustomerID → CustomerName, CustomerAddress

ProductID → ProductDescription, ProductFinish, ProductStandardPrice

OrderID, ProductID → OrderQuantity

Therefore, NOT in 2nd Normal Form

Removing partial dependencies



**Getting it into
Second Normal
Form**

Partial dependencies are removed, but there are still transitive dependencies

THIRD NORMAL FORM

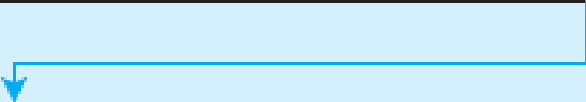
- 2NF PLUS ***no transitive dependencies*** (functional dependencies on non-primary-key attributes)
- Note: This is called transitive, because the primary key is a determinant for another attribute, which in turn is a determinant for a third
- ***Solution:*** Non-key determinant with transitive dependencies go into a new table; non-key determinant becomes primary key in the new table and stays as foreign key in the old table

Removing partial dependencies

<u>OrderID</u>	OrderDate	<u>CustomerID</u>
----------------	-----------	-------------------

ORDER (3NF)

Getting it into
Third Normal
Form



<u>CustomerID</u>	CustomerName	CustomerAddress
-------------------	--------------	-----------------

CUSTOMER (3NF)

Transitive dependencies are removed.